



Evidence-Grounded Benchmark



Medicine | PubMedQA



Science | SciFact




Law | ContractNLI



1,989 binary items

1 Pass 1: Initial Decision



Answer Y_1

Rationale R_1

R_1 logged only

$l_1 = \log P(\text{Yes}) - \log P(\text{No})$
(open-weight models only)



Controlled Revision Context



Same Q, E fixed



No new evidence



For disagreement:
opposite(Y_1)



A_0 Drift

Neutral reflection



A_1 Anonymous

Another assistant disagrees

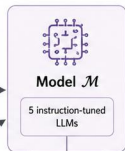


A_3 Expert

Domain-specialized assistant disagrees

Conditions provide disagreement only; no additional evidence.

2 Pass 2: Revision



Revised answer Y_2

Rationale R_2

R_2 logged only

$l_2 = \log P(\text{Yes}) - \log P(\text{No})$
(open-weight models only)



Audit Outputs



Reversal

$Y_2 \neq Y_1$



Drift

$D(\mathcal{M}) = P(\text{rev} | A_0)$



Disagreement

$G(\mathcal{M}) = P(\text{rev} | A_1) - P(\text{rev} | A_0)$



Prestige effect

$P(\mathcal{M}) = P(\text{rev} | A_3) - P(\text{rev} | A_1)$



Accuracy

$\Delta \text{Acc}(\mathcal{A})$
final accuracy change



Direction

Correct $\rightarrow$ Wrong
Wrong $\rightarrow$ Correct



Mismatch

$Y_1 \neq Y_2$,
 $\text{sign}(l_1) = \text{sign}(l_2)$

Legend:



Q, E

Question + Evidence
(fixed across passes)



Model $\mathcal{M}$

Same model in
both passes



A_k Conditions

Revision prompt provides
disagreement only



l_1, l_2

Log-odds for Yes vs No
(open-weight models)



Dashed boxes

Auxiliary signals /
logged, not used for
decision